

# Federated Learning Experiment Summary: MNIST Classification with FedAvg

July 23, 2025

## Abstract

This document summarizes a federated learning (FL) experiment for MNIST digit classification using the Flower framework with Federated Averaging (FedAvg). The setup involves three clients training a neural network on non-IID MNIST data partitions, achieving a global accuracy of 97.58% after 20 rounds. The process, performance, limitations, and recommendations for improvement are detailed, based on provided code, configurations, and logs.

## 1 Experiment Process

### 1.1 Setup and Configuration

The experiment trains a feedforward neural network (MNISTNet,  $784 \rightarrow 128 \rightarrow 64 \rightarrow 10$ ,  $\sim 109,386$  parameters) on the MNIST dataset using Federated Averaging (FedAvg) within the Flower framework. Key configurations include:

- Clients: 3
- Rounds: 20
- Local Epochs: 50
- Batch Size: 32
- Learning Rate: 0.001 (SGD optimizer)
- Non-IID Alpha: 0.1 (highly skewed data)
- Device: CPU

The MNISTDataPartitioner (in `data.py`) splits MNIST into non-IID partitions, with each client having  $\sim 20,000$  samples, 80% from specific digits (e.g., Client 0: digits 0–3, Client 1: digits 4–6, Client 2: digits 7–9) and 20% from others, controlled by `non_iid_alpha=0.1`.

### 1.2 FedAvg Technique

FedAvg involves:

- Clients training local models for 50 epochs using SGD.
- Sending updated parameters to the server.

- Server averaging parameters weighted by sample count:

$$w_{t+1} = \sum_{k=1}^K \frac{n_k}{n} w_{t+1}^k$$

where  $w_{t+1}^k$  is client  $k$ 's updated model,  $n_k$  is its sample count, and  $n$  is the total samples.

The server evaluates the global model on a centralized test set (10,000 samples) after each round.

### 1.3 Execution

The experiment is orchestrated by `main.py`, which:

- Loads and partitions data using `MNISTDataPartitioner`.
- Initializes clients via `client_fn` (in `main.py`) and the server via `MNISTFedAvgStrategy` (in `server.py`).
- Runs the simulation using Flower's `start_simulation`.

Results are saved to `logs/experiment_20250722_025634`. Client 0 reports a local accuracy of 98.74% and loss of 0.0386 after 50 epochs.

### 1.4 Initial Failure and Fix

Early experiments failed (accuracy  $\sim 12.89\%$ , loss  $\sim 2.3059$ ), matching random model performance, due to a bug in `client.py`'s `client_fn`, which set `train_dataloader=None`, preventing training. This was fixed by using the correct `client_fn` in `main.py`, which assigns valid dataloaders from `CLIENT_DATALOADERS`.

## 2 Performance

The global model improves significantly:

- Round 0: Loss = 2.4658, Accuracy = 12.32% (random).
- Round 5: Loss = 0.1101, Accuracy = 96.65%.
- Round 10: Loss = 0.0894, Accuracy = 97.30%.
- Round 15: Loss = 0.0849, Accuracy = 97.52%.
- Round 20: Loss = 0.0842, Accuracy = 97.58%.

Client 0 achieves 98.74% accuracy and 0.0386 loss locally. The high performance despite non-IID data (`non_iid_alpha=0.1`) is due to:

- Extended Training: 20 rounds and 50 epochs allow robust local convergence.
- FedAvg: Aggregates diverse client models effectively.
- Data Diversity:  $\sim 20,000$  samples per client with some non-dominant classes.
- Model Capacity: MNISTNet is well-suited for MNIST.

Slight loss fluctuations in later rounds (e.g., 0.0837 in Round 19 to 0.0842 in Round 20) reflect non-IID data challenges.

### 3 Limitations

1. Deprecated Flower Features:
  - `start_simulation` is deprecated; Flower recommends the `flwr run` CLI.
  - `client_fn` returns `NumPyClient` instead of `Client`, causing a warning.
2. No Client-Side Evaluation: `fraction_evaluate=0.0` skips local validation, limiting insights into client generalization.
3. No Metrics Aggregation: Missing `fit_metrics_aggregation_fn` prevents aggregation of client training metrics.
4. Conservative Optimizer: SGD with learning rate 0.001 is slow; Adam could converge faster.
5. No Early Stopping: Accuracy plateaus around 97.5% after Round 15, wasting computation.
6. Incorrect Summary Output: `print_final_results` in `main.py` previously reported invalid metrics (e.g., 230.59% accuracy).

### 4 Recommendations

1. Adopt Flower CLI Workflow:
  - Use `flwr new` and `flwr run` to create and run a Flower app, restructuring code into a `pyproject.toml`-based project (see <https://flower.ai/docs/framework/how-to-run-simulations.html>).
2. Fix `client_fn` Warning:
  - Update `client_fn` in `main.py` to return `Client` using `to_client()`:

```
from flwr.common import Context
import flwr as fl
def client_fn(context: Context) -> fl.client.Client:
    client_id = int(context.node_config.get("partition-id", context.node_id))
    train_loader = CLIENT_DATALOADERS[client_id]['train']
    val_loader = CLIENT_DATALOADERS[client_id]['val']
    client = create_client(client_id=client_id, train_dataloader=train_loader,
                           val_dataloader=val_loader, epochs=EXPERIMENT_CONFIG['local_epochs'],
                           learning_rate=EXPERIMENT_CONFIG['learning_rate'],
                           device=EXPERIMENT_CONFIG['device'])
    return client.to_client()
```
3. Enable Client-Side Evaluation:
  - Set `fraction_evaluate=0.5` and `min_evaluate_clients=2` in `server.py`'s `create_server_strategy` to evaluate client models.
4. Aggregate Client Metrics:
  - Add `fit_metrics_aggregation_fn` in `server.py` to compute weighted average of client accuracies:

```
def fit_metrics_aggregation_fn(fit_metrics):
    total_samples = sum(metrics["num_samples"] for _, metrics in fit_metrics)
    weighted_acc = sum(metrics["accuracy"] * metrics["num_samples"]
                        for _, metrics in fit_metrics) / total_samples
    return {"accuracy": weighted_acc}
```

- Update MNISTFlowerClient.fit to include num\_samples in metrics.

#### 5. Use Adam Optimizer:

- Update train\_model in utils.py to use Adam with learning rate 0.01:

```
def train_model(model, dataloader, epochs=50, learning_rate=0.01, device="cpu"):
    model.to(device)
    model.train()
    optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
    criterion = torch.nn.CrossEntropyLoss()
    # ... (rest unchanged)
```

- Update main.py configuration: learning\_rate=0.01.

#### 6. Add Early Stopping:

- Modify create\_evaluate\_fn in server.py to stop if accuracy improvement is small (e.g., <0.01% after Round 5 and accuracy >97%).

#### 7. Fix Summary Output:

- Update print\_final\_results in main.py to correctly parse metrics:

```
def print_final_results(history):
    print("\n" + "="*60 + "\nFEDERATED LEARNING EXPERIMENT RESULTS\n" + "="*60)
    if hasattr(history, 'metrics_centralized') and history.metrics_centralized.get('accuracy'):
        print("\n Global Model Performance:")
        for round_num, acc in history.metrics_centralized['accuracy']:
            loss = next(l for r, l in history.losses_centralized if r == round_num)
            print(f" Round {round_num}: Loss = {loss:.4f}, Accuracy = {acc:.2f}%")
```

## 5 Expected Outcome

With these improvements:

- Global accuracy may exceed 98% with faster convergence using Adam.
- Client-side evaluation will provide insights into local model performance.
- Early stopping will reduce computation (e.g., stopping around Round 15).
- Compatibility with future Flower versions will be ensured.